

# CompanyMind

---

## AI-Powered Enterprise Knowledge Base

---

### Semantic Search & Retrieval-Augmented Generation

---

MongoDB Atlas Vector Search · React 18 · Groq Llama 3 70B · FastAPI · MiniLM-L6-v2

*"Stop searching for keywords. Start finding answers."*

<b>Repository</b>	<a href="https://github.com/nikhilkumarpanigrahi/Companymind">github.com/nikhilkumarpanigrahi/Companymind</a>
<b>License</b>	Apache 2.0
<b>Date</b>	February 2025

# 1. Problem & Solution

---

## The Problem

Traditional enterprise search fails at three levels:

- **Keyword mismatch** — Searching "how to speed up database queries" misses a document titled "Database Indexing Strategies" because exact words don't match
- **No semantic understanding** — keyword search can't bridge vocabulary gaps between question and answer
- **No synthesis** — even when documents are found, users must read and piece together answers manually

Organizations lose **20% of productive time** searching for information they already have (McKinsey).

## Our Solution


CompanyMind converts every document into a **384-dimensional vector embedding** that captures semantic meaning. When you search, your query is embedded into the same vector space and MongoDB Atlas finds the most similar documents using **cosine similarity** — understanding intent, not matching strings.


The **RAG pipeline** goes further: it retrieves top matching documents and feeds them to **Llama 3 70B**, which generates a comprehensive, cited answer — streamed token-by-token via **Server-Sent Events**.

Mode	How It Works	Latency
<b>Semantic Search</b>	Query → 384-dim embedding → MongoDB <code>\$vectorSearch</code> → ranked results with cosine scores	< 200ms
<b>Ask AI (RAG)</b>	Query → retrieve top 5 docs → assemble context → Llama 3 70B generates cited answer → SSE stream	1–3s

## 2. Architecture

---





## Tech Stack

Layer	Technology	Purpose
<b>Frontend</b>	React 18 + TypeScript + Vite 6 + Tailwind CSS 3	SPA with dark theme, SSE streaming, voice input
<b>Backend</b>	Express.js 4 + Zod + Helmet + compression	REST API with validation, security, rate limiting
<b>Database</b>	MongoDB Atlas	<code>\$vectorSearch</code> (HNSW, cosine, 384-dim) + <code>\$text</code> (BM25)
<b>Embeddings</b>	FastAPI + SentenceTransformers <code>all-MiniLM-L6-v2</code>	384-dim L2-normalized vectors, self-hosted, zero API costs
<b>LLM</b>	Groq API + Llama 3.3 70B Versatile	RAG answer generation with real-time SSE streaming
<b>Deployment</b>	Docker multi-stage, Render, AWS ECS Fargate	Non-root containers, health checks, zero-download images

## 3. Core Innovation: Hybrid Search with RRF

CompanyMind's key technical contribution is its **Reciprocal Rank Fusion (RRF)** hybrid search — combining vector similarity and BM25 keyword relevance into a single, parameter-free ranking.

### Why Not Simple Score Interpolation?

Linear interpolation ( `$\alpha \cdot \text{vector\_score} + (1-\alpha) \cdot \text{text\_score}$` ) requires careful tuning of  $\alpha$ , score normalization, and assumes comparable distributions. **RRF uses only rank positions**, making it robust regardless of score scales.

### Algorithm

```
function reciprocalRankFusion(vectorResults, textResults, k = 60) {
  const scoreMap = new Map();

  vectorResults.forEach((doc, rank) => {
    const id = doc._id.toString();
    scoreMap.set(id, (scoreMap.get(id) || 0) + 1 / (k + rank + 1));
  });

  textResults.forEach((doc, rank) => {
    const id = doc._id.toString();
    scoreMap.set(id, (scoreMap.get(id) || 0) + 1 / (k + rank + 1));
  });

  return sorted(scoreMap); // Documents in BOTH lists get boosted scores
}
```

Documents in **both** result sets receive higher combined scores. Threshold filtering ( `$1/(K + \text{limit}^3)$` ) removes noise. K=60 follows the [original RRF paper](#).

### 4 Search Strategies Compared

Strategy	MongoDB Operator	What It Measures
<b>Regex</b>	<code>\$regex</code>	Baseline — brute-force pattern match
<b>Full-Text</b>	<code>\$text</code> + textScore	Traditional keyword search with stemming
<b>Vector</b>	<code>\$vectorSearch</code> cosine	Pure semantic similarity (384-dim HNSW)
<b>Hybrid RRF</b>	Vector + Text + RRF merge	Best-of-both-worlds fusion

The built-in **benchmark suite** runs all 4 strategies on the same query, reporting per-method latency, result overlap matrix, and average relevance scores — quantitatively proving hybrid outperforms any single method.

## 4. RAG Pipeline — How Ask AI Works

```
User Question: "How does vector search find relevant documents?"
|
1. EMBED QUERY -----> MiniLM-L6-v2 -> 384-dim vector (~50ms, cached <1ms)
|
2. VECTOR SEARCH -----> MongoDB $vectorSearch, HNSW, top 5 results (~40ms)
|
3. CONTEXT ASSEMBLY -----> Each source: title + relevance % + content (1500
chars)
|
|
4. LLM STREAMING -----> Groq Llama 3 70B, temp=0.3, max 1024 tokens
|
|
5. CITED ANSWER -----> Every claim backed by [Source: title] references
```

### What Makes the RAG Unique

- **Context-grounded** — System prompt forces LLM to use *only* retrieved documents, preventing hallucination
- **Multi-turn memory** — Sliding window of last 3 Q&A pairs enables natural follow-up questions
- **Nginx-aware SSE** — **X-Accel-Buffering: no** header prevents proxy buffering; errors sent as SSE events, never dropping the connection
- **Source transparency** — Sources displayed *immediately* before the answer starts streaming

## 5. Triple-Layer Caching

CompanyMind implements **3 independent LRU caches** that progressively eliminate latency:

```

Request: "vector databases"
|
[Tier 3] Search Result Cache -----> 200 entries, 5-min TTL
| miss
| Key: "hybrid:<query>:<limit>"
|
[Tier 2] Node Embedding Cache -----> 500 entries, 10-min TTL
| miss
| Key: text.trim().toLowerCase()
|
[Tier 1] Python Embedding Cache --> 2,000 entries, no TTL
| miss
| Key: MD5(text)
|
MinILM-L6-v2 Inference (~50ms)

```

Tier	Location	Max Size	TTL	Cache-Hit Latency
Search results	Node.js	200	5 min	< 1ms
Embeddings	Node.js	500	10 min	< 1ms
Embeddings	Python	2,000	None	< 1ms (skips model)

On a cache-warm system, repeated queries return in **< 1ms**.

## 6. Embedding Model: Why MiniLM-L6-v2

Property	Value	Advantage
Architecture	6-layer Transformer	2x faster than 12-layer BERT
Dimensions	384	50% less storage than BERT (768), 75% less than OpenAI (1536)
Normalization	L2-normalized	Cosine similarity = dot product (MongoDB optimizes this)
Model size	~80 MB	Fits in Docker image, no GPU needed
Inference	~50ms/query (CPU)	Fast enough for real-time search
Hosting	Self-hosted (FastAPI)	Zero API costs, no rate limits, full data privacy

The embedding engine uses a **singleton pattern** — model loaded once at startup, reused for all requests, avoiding the ~2s cold-start per request.

## 7. API Overview

---

### Core Endpoints (Express.js — Port 8080)

Method	Endpoint	Purpose
GET	/health	Server health + DB status (always 200)
POST	/api/documents	Create document (auto-embedded)
GET	/api/documents	List documents
GET	/api/documents/stats	Category/tag analytics
GET	/api/search?q=...	<b>Hybrid RRF search</b> (primary)
POST	/api/search	Pure vector search (API-to-API)
POST	/api/ask	RAG answer (non-streaming)
POST	/api/ask/stream	<b>RAG answer (SSE streaming)</b>
GET	/api/ask/analytics	Query metrics dashboard
POST	/api/benchmark	Run 4-strategy benchmark

### Embedding Service (FastAPI — Port 8000)

Method	Endpoint	Purpose
GET	/health	Model status + cache hit rate
POST	/embed-query	Single text → 384-dim vector
POST	/embed-batch	Batch texts → vectors (max 512)

### Validation & Security

- **Zod schemas** on every endpoint — title (1–300), content (1–10,000), query (1–1,000), question (1–2,000)
- **Rate limiting** — 100 req/min general, 30 req/min search/ask
- **Helmet** security headers, **1 MB body limit**, **CORS**, **non-root Docker**

## 8. Performance Metrics

Metric	Value	Details
Embedding generation	~50ms	MiniLM-L6-v2, 384 dimensions, CPU-only
Vector search	20–80ms	MongoDB Atlas <code>\$vectorSearch</code> HNSW index
Hybrid search (end-to-end)	< 200ms	Including embedding + parallel search + RRF merge
RAG answer	1–3s	Groq Llama 3 70B, 1024 max tokens, SSE stream
Cache hit (any tier)	< 1ms	3-tier LRU eliminates redundant computation

### Example Semantic Queries

Query	What It Finds (No Keyword Overlap!)
"speed up my app"	Caching, indexing, optimization docs
"protecting against cyber threats"	XSS/CSRF, OWASP, DevSecOps docs
"how machines learn from data"	ML, deep learning, neural network docs
"how transformers work in AI"	Attention mechanisms, BERT, transformer architecture

## 9. Design Decisions & Trade-offs

Decision	Rationale
<b>Separate embedding microservice</b>	ML model in Python (best ecosystem), API in Node.js (best for Express). Language-agnostic scaling.
<b>In-memory LRU over Redis</b>	Zero infrastructure. For finite knowledge base queries, LRU provides sub-ms lookups without ops burden.
<b>CPU-only PyTorch</b>	MiniLM runs in ~50ms on CPU. GPU would add ~3 GB to Docker image + require NVIDIA runtime.
<b>RRF over learned ranking</b>	Parameter-free, robust. Learning-to-rank needs training data — overkill for this scope.
<b>Groq over OpenAI</b>	Groq's Llama 3 70B in 1–3s enables real-time SSE streaming. Free tier is generous.
<b>CommonJS (.cjs) backend</b>	Mongoose + some deps most reliable with <code>require()</code> . ESM still inconsistent server-side.
<b>\$vectorSearch over custom HNSW</b>	Atlas handles index maintenance, replication, scaling. No reinventing the wheel.

**Resilient startup**

Server listens *before* DB connects. Health checks always respond. Critical for Render/ECS cold starts.

## 10. What Makes CompanyMind Different

---

1. **Hybrid RRF search** — mathematically proven fusion of semantic + keyword relevance, not just one or the other
  2. **Triple-layer caching** — three independent LRU caches reduce repeat query latency to < 1ms
  3. **Production-ready SSE streaming** — Nginx-aware, graceful error events, never drops connection
  4. **Multi-turn RAG conversations** — sliding window context for natural follow-ups
  5. **Built-in benchmark suite** — quantitatively proves vector search superiority with overlap analysis
  6. **Zero-download Docker** — ML model baked into image via multi-stage build
  7. **Voice search** — Web Speech API for hands-free knowledge retrieval
  8. **Self-hosted embeddings** — zero API costs, no rate limits, full data privacy
- 

**CompanyMind** — Built for the MongoDB AI Hackathon

*Turning documents into knowledge, and knowledge into answers.*

Apache License 2.0 · [github.com/nikhilkumarpanigrahi/Companymind](https://github.com/nikhilkumarpanigrahi/Companymind)